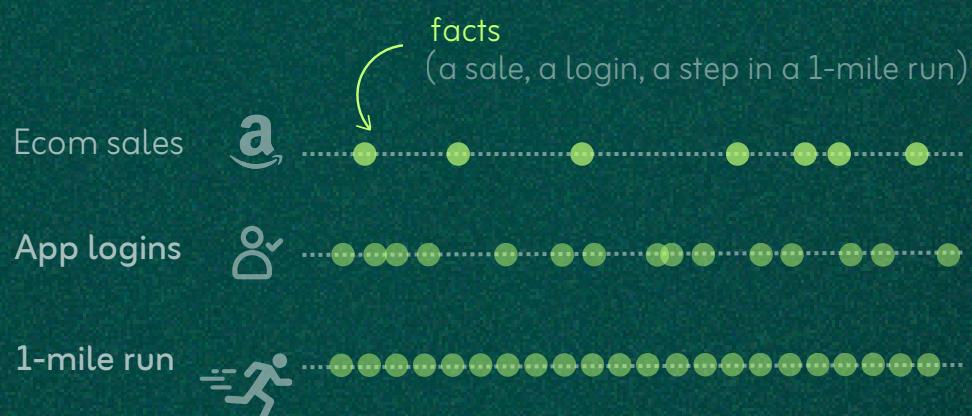


## 1 FACT ~ Truth

- Atomic (can't be broken down into smaller units)
- Can be aggregated up/down
- Not slowly changing
- Can't be changed back



## 2 Why Fact Data is hard to model?

### 1. Scalability (10-100x the amount of dimension data)

Q: How many actions taken per dimension?

Ie. : 2 bio users x 30 notifications/day ~ 50 bio notifications/day

### 2. Need more context for effective analysis



Q: What data dimensions do we need to make the fact data more valuable?

Ie. A user's location bring context (what's the clickthrough rate between US and Canada notification recipients?)

### 3. Duplicates are more common. Causes:

- Bugs  
Ie. A software engineer pushed a bug in the logger .....> 1 click generates 2 logs .....> Data Quality issue
- Genuine duplicates (can be handled via deduplication)  
Ie. A notification clicked now and clicked again 3 hours later  
Tracking 2 genuine actions without counting them as 2? .....> deduplication  
No deduplication .....> your metric will look weird! (ie. a 200% CTR)

## 4 How Fact Data Modeling works?

### Normalization

- Splits data across multiple tables
- Minimizes data redundancy
- Removes duplicates, increases data integrity

### Denormalization

- Combines data into a single table
- Faster querying (use GROUP BY, not JOIN)
- High storage cost, potential duplications

Trade-off



### Fact Data vs. Raw Logs: same thing?

- No, but both interconnected (no fact without raw logs)
- The level of trust of fact data should be higher than the trust in raw logs

#### Raw Logs

- Data for systems, server running.
- Usually owned by SEs (software engineers)
- Ugly schemas, potential duplicates, DQ issues
- Short retention

#### Fact Data

- Data for analysis.
- Usually owned by DEs (data engineers)
- Nice schemas/column names
- Quality guarantees (uniqueness, not nulls, ...)
- Longer retention

 Goal as a DE: collaborate closely with SEs, be able to convert raw logs into highly trusted fact data

### Aspects of Fact Modeling

#### A WHO ..... → ‘Who is part of that event?’

- ‘Pushed out as IDs (ie. Tesla car ..... → <‘user\_id’, ‘car\_id’, ‘operating\_system\_id’>)

#### B WHERE ..... → ‘Where is the event taking place?’

- Modeled with ids (Who) + additional dimensions  
ie. location (Country, City, State); app pages (Homepage, Profile Page... modeled w/ backslashes)

#### C HOW ..... → ‘How is part of the event?’

- Similar to ‘Where’ and ‘Who’  
ie. “Make a click using an iPhone”
- where (where in the app, which button they clicked)  
→ how (methodology that lead to click)

#### D WHAT ..... → ‘What is the event taking place?’

- Atomic events  
ie. 1. Atomic : A notification ..... → ‘generated’, ‘sent’, ‘clicked’, ‘converted into purchase’  
2. Not Atomic: ‘Burning Man’ ..... → Higher-level event (group/aggregation of smaller events)

#### E WHEN ..... → ‘When is the event taking place?’

- Modeled as timestamp
- Align all app logins to same time zone (UTC), client-side login (higher fidelity vs server side login)



## FACT DATA MODELING FUNDAMENTALS (II)

### MUSTs of FACT data

- Quality guarantees: no duplicates, no NULLs in ‘what’, ‘when’ event fields
- Size: keep it small, FACT < RL\* (what’s needed for analysis only )
- Parse out hard-to-understand columns

- ✓✓ Strings, integers, decimals, enumerates  
 ✓ Array of strings  
 ✗ JSON files

### WHEN to model dimensions in FACT data



#### Network Logs (Traffic) Pipeline

- Goal: understand how Netflix microservices (apps) ‘talk’ to each other. Analyze traffic data  
ie. What is the impact of app A ‘talking’ to n apps if being hacked?
- Problem: Huge dataset (2PB/day)
- Solution (2-sided):

##### 💡 IPv4 IP addresses (32-bit)

- IP addresses as app identifier, stored in a smaller database!
- BROADCAST JOIN (Spark) of IDs with Network Logs data

- Learning for Data Engineers:

✓ DO: Go upstream & solve the problem at source.

→ 1028x the IPv4 volume!

##### 💡 IPv6 IP addresses (128-bit hexadecimal)

Denormalize FACT data & log it ahead of time.  
To get rid of the join, Each app adopts a “sidecar proxy” to log of which app they are

✗ DONT: Optimize a pipeline at first

### HOW logging fits into FACT data

- Brings in all critical context (columns)  
Collaborative work with system engineers (handle event generation in the app)
- Necessary logs only (RL ~ expensive)
- Comformance: logging needs to be in some type of contract, shared schema/vision



#### Setting Pricing Comformance at

- Goal: What describes a price, what pricing is available?
- Problems: two teams (Pricing team Online Service Team), two programming frameworks (Ruby, Scala), potential difference in pricing understanding/definition.  
How to ensure definition and replication on both framework is aligned?
- Solution: shared schema (middle layer)

##### Thrift Schema

- Way to describe schemas/functions in a specific, language agnostic way  
Ruby/Scala reference both to same schema for price (ie. Scala adds a new pricing column, then schema sends push notification to Ruby to update code)

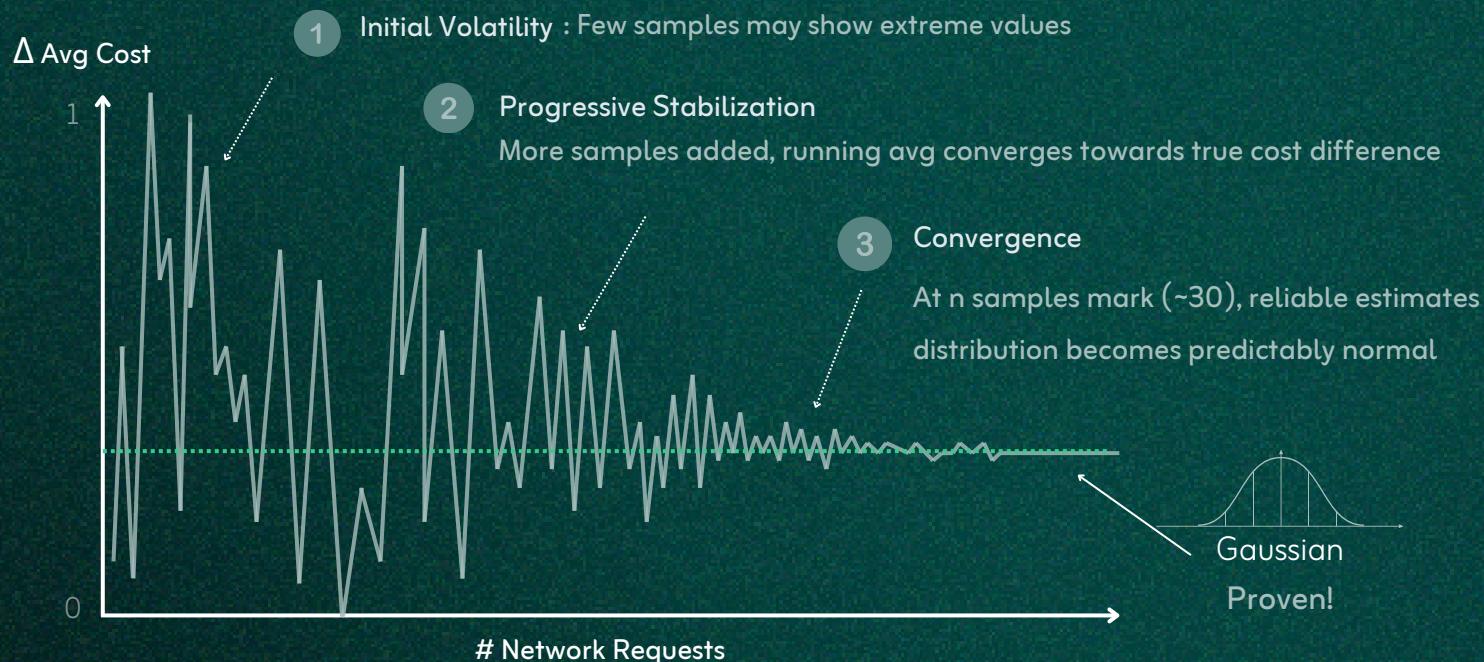


**FACT DATA MODELING FUNDAMENTALS (II)****1 SAMPLING** ~ small data subset; validated by law of Large Numbers

- Works for problems/ metrics that gauge directionality of things.
- Does **not** work for security related problems, low probability events, specific row data events.

**The 'Infrastructure Cost' metric case at NETFLIX**

Q: Does an A/B test cause a higher AWS S3/EC2 infra cost in test group vs control group?

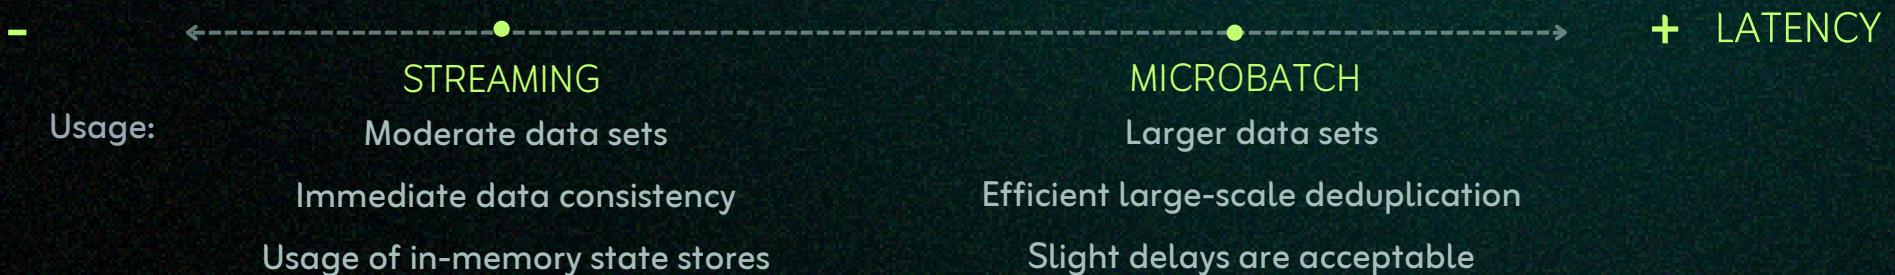
**1 BUCKETING**

- Bucket on the 'who' (ie. `user_id`)
- Join the bucketed column w/o shuffling across the entire dataset. Faster, help minimize shuffle in Spark. SMB join (sorted merged bucket)  $\rightarrow$  joins w/o shuffle at all. Both buckets lined up & sorted

**2 How Long to hold onto FACT data?****3 How to handle duplicates in data?**

**Deduplication** ~ remove duplicate/redundant copies of data (storage & efficiency improvement)

- Analyze distribution of duplicates over time periods (week/day/hour) & set a timeframe where deduplication matters (latency).



## B MICROBATCH ~ Example: Hourly Deduplication with daily microbatch

